

LearnTrack Big Data Project

Complete Setup, Python Scripts Execution & Visualization Guide

2025–2026

Contents

1	Project Overview	2
2	Docker Network Creation	2
3	MongoDB Setup	2
3.1	Run MongoDB Container	2
3.2	Insert Sample Data	2
4	Redis Setup (Python Script Execution)	3
4.1	Run Redis Container	3
4.2	Execute Redis Python Loader	3
4.3	Verify Redis Scores	3
5	Neo4j Setup (Python Script Execution)	3
5.1	Run Neo4j Container	3
5.2	Execute Neo4j Python Loader	3
5.3	Verify Graph in Neo4j Browser	3
5.4	Multi-Region Access Detection	4
6	Cassandra Setup	4
6.1	Run Cassandra Container	4
6.2	Create Keyspace and Table	4
7	Apache Spark Batch Processing	4
7.1	Run Spark Container	4
7.2	Execute Spark Job	5
7.3	Verify Cassandra Results	5
8	Grafana Visualization	5
8.1	Run Grafana Container	5
8.2	Access Grafana	5
8.3	Dashboards	5
9	Conclusion	5

1 Project Overview

LearnTrack is a polyglot Big Data analytics platform designed to analyze student activities in online learning systems.

The platform provides:

- Real-time engagement scoring
- Multi-region access detection
- Batch analytics at course level
- Interactive dashboards

Technologies used: MongoDB, Redis, Neo4j, Apache Spark, Cassandra, and Grafana.

2 Docker Network Creation

All containers communicate through a shared Docker network.

```
1 docker network create learntrack-net
```

3 MongoDB Setup

3.1 Run MongoDB Container

```
1 docker run -d \
2   --name mongo \
3   --network learntrack-net \
4   -p 27017:27017 \
5   mongo
```

3.2 Insert Sample Data

```
1 docker exec -it mongo mongosh
```

```
1 use learntrack
2
3 db.activities.insertMany([
4   { studentId:"S7", courseId:"C1", action:"page_view", timestamp:new
5     Date(), region:"AF" },
6   { studentId:"S7", courseId:"C1", action:"page_view", timestamp:new
7     Date(), region:"EU" },
8   { studentId:"S9", courseId:"C3", action:"quiz_attempt", timestamp:new
9     Date(), region:"US" }
])
10
11 db.activities.find().pretty()
```

4 Redis Setup (Python Script Execution)

4.1 Run Redis Container

```
1 docker run -d \
2   --name redis \
3   --network learntrack-net \
4   -p 6379:6379 \
5   redis
```

4.2 Execute Redis Python Loader

Redis engagement scores are computed using the Python script `redis_loader.py`, which reads activity data from MongoDB and stores aggregated scores in Redis.

```
1 docker exec -it spark bash
1 python /opt/spark/work-dir/redis_loader.py
```

4.3 Verify Redis Scores

```
1 KEYS student:*:score
2 GET student:S7:score
3 MGET student:S7:score student:S9:score
```

5 Neo4j Setup (Python Script Execution)

5.1 Run Neo4j Container

```
1 docker run -d \
2   --name neo4j \
3   --network learntrack-net \
4   -p 7474:7474 \
5   -p 7687:7687 \
6   -e NEO4J_AUTH=neo4j/password \
7   neo4j:5
```

5.2 Execute Neo4j Python Loader

The graph database is populated using the Python script `neo4j_loader.py`, which loads activities from MongoDB into Neo4j.

```
1 docker exec -it spark bash
1 python /opt/spark/work-dir/neo4j_loader.py
```

5.3 Verify Graph in Neo4j Browser

Open:

<http://localhost:7474>

```
1 MATCH (s:Student)-[:PERFORMED]->(a:Activity)-[:IN_REGION]->(r:Region)
2 RETURN s, a, r;
```

5.4 Multi-Region Access Detection

```
1 MATCH (s:Student)-[:PERFORMED]->(a1)-[:IN_REGION]->(r1),  
2     (s)-[:PERFORMED]->(a2)-[:IN_REGION]->(r2)  
3 WHERE r1.name <> r2.name  
4     AND duration.between(a1.timestamp, a2.timestamp)  
5         < duration({minutes:30})  
6 RETURN s.id AS studentId,  
7         collect(DISTINCT r1.name) + collect(DISTINCT r2.name) AS regions  
8 ;
```

6 Cassandra Setup

6.1 Run Cassandra Container

```
1 docker run -d \  
2   --name cassandra \  
3   --network learntrack-net \  
4   -p 9042:9042 \  
5   cassandra:4
```

6.2 Create Keyspace and Table

```
1 docker exec -it cassandra cqlsh
```

```
1 CREATE KEYSPACE learntrack  
2 WITH replication = {'class':'SimpleStrategy','replication_factor':1};  
3  
4 USE learntrack;  
5  
6 CREATE TABLE course_engagement (  
7   courseid text,  
8   date date,  
9   count int,  
10  PRIMARY KEY (courseid, date)  
11 );
```

7 Apache Spark Batch Processing

7.1 Run Spark Container

```
1 docker run -it \  
2   --name spark \  
3   --network learntrack-net \  
4   -p 4040:4040 \  
5   bitnami/spark:3 \  
6   bash
```

7.2 Execute Spark Job

The Spark job reads activities from MongoDB and writes aggregated engagement results into Cassandra.

```
1 spark-submit \
2   --packages \
3     org.mongodb.spark:mongo-spark-connector_2.12:10.3.0, \
4     com.datastax.spark:spark-cassandra-connector_2.12:3.5.0 \
5     /opt/spark/work-dir/spark_job.py
```

7.3 Verify Cassandra Results

```
1 SELECT * FROM learntrack.course_engagement;
```

8 Grafana Visualization

8.1 Run Grafana Container

```
1 docker run -d \
2   --name grafana \
3   --network learntrack-net \
4   -p 3000:3000 \
5   grafana/grafana
```

8.2 Access Grafana

<http://localhost:3000>

- Username: admin
- Password: admin

8.3 Dashboards

- Redis: student engagement scores
- Cassandra: course-level engagement analytics
- Neo4j: multi-region access detection

9 Conclusion

This document presented a complete and consistent setup of the LearnTrack Big Data platform, including Docker orchestration, Python-based data loading, batch analytics with Spark, and real-time visualization using Grafana.